Production ML Pipelines

Production ML Pipelines

01	Ways to do ML on Google Cloud
02	Vertex AI Pipelines
03	AI Hub

Production ML Pipelines

01	Ways to do ML on Google Cloud
02	Vertex AI Pipelines
03	AI Hub

Create and deploy custom models with Vertex AI

Cloud TPUs

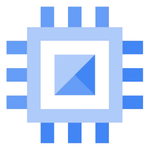

Compute Engine

Dataproc

Google Kubernetes Engine

Vertex AI

BigQuery ML

AutoML

Cloud Translation API

Vision API

Speech-to-Text API

Video Intelligence API

Data Loss Prevention API

Text-to-Speech API

Cloud Natural Language API

Dialogflow

Build a Custom Model

Build Custom Model (codeless)

Call a Pretrained Model

Vertex AI is a fully managed service for custom machine learning models

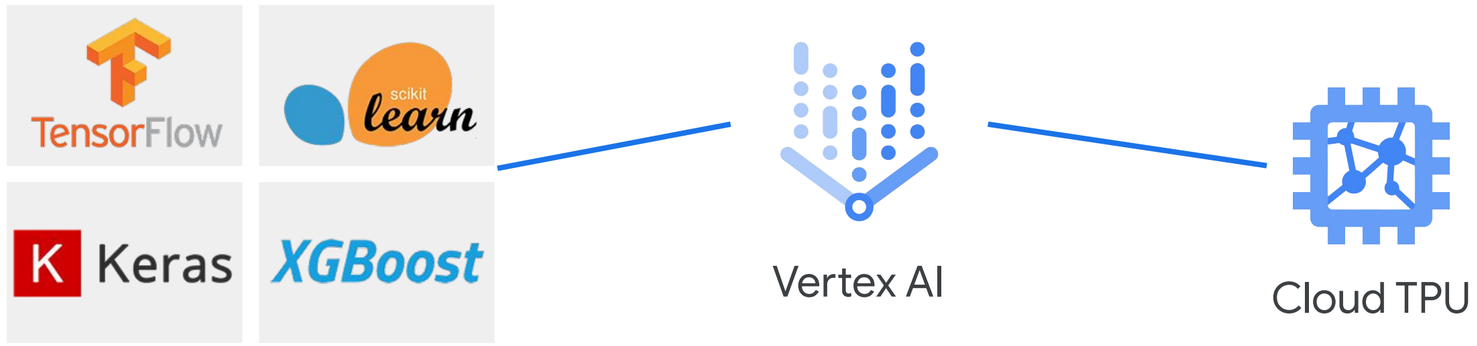

- Scales to production
- Batching and distribution of model training
- Performs transformations on input data
- Hyper-parameter tuning
- Host and autoscale predictions
- Serverless - self-tuning - manages overhead

In this course, we don't cover writing TensorFlow models, only ways to operationalize them

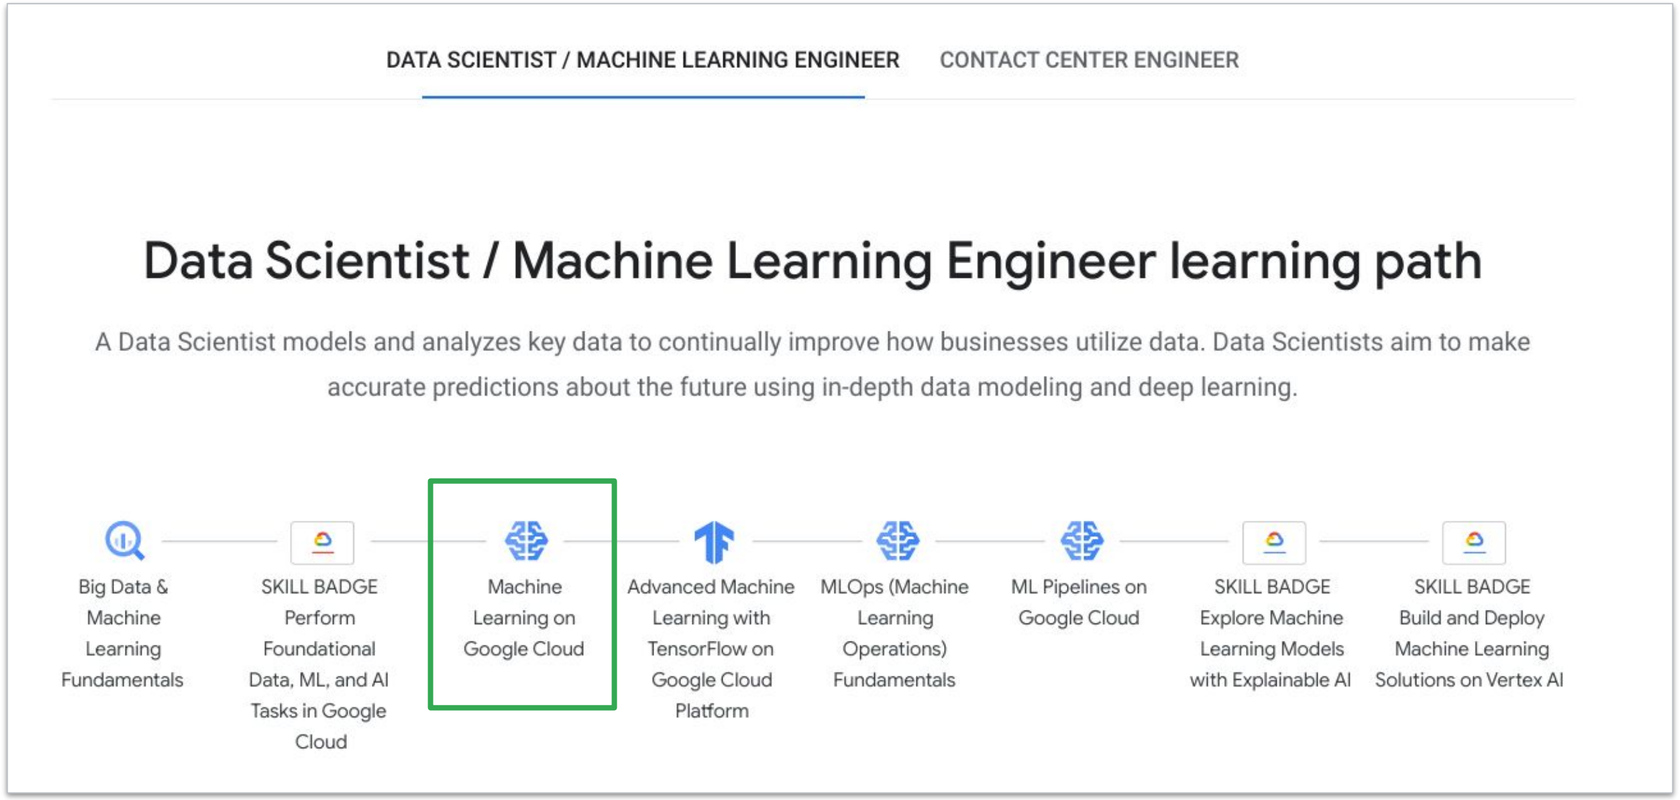

Google Cloud Training - Machine Learning and AI

Production ML Pipelines

01	Ways to do ML on Google Cloud
02	Vertex AI Pipelines
03	AI Hub

Pipelines automate the training and deployment of models

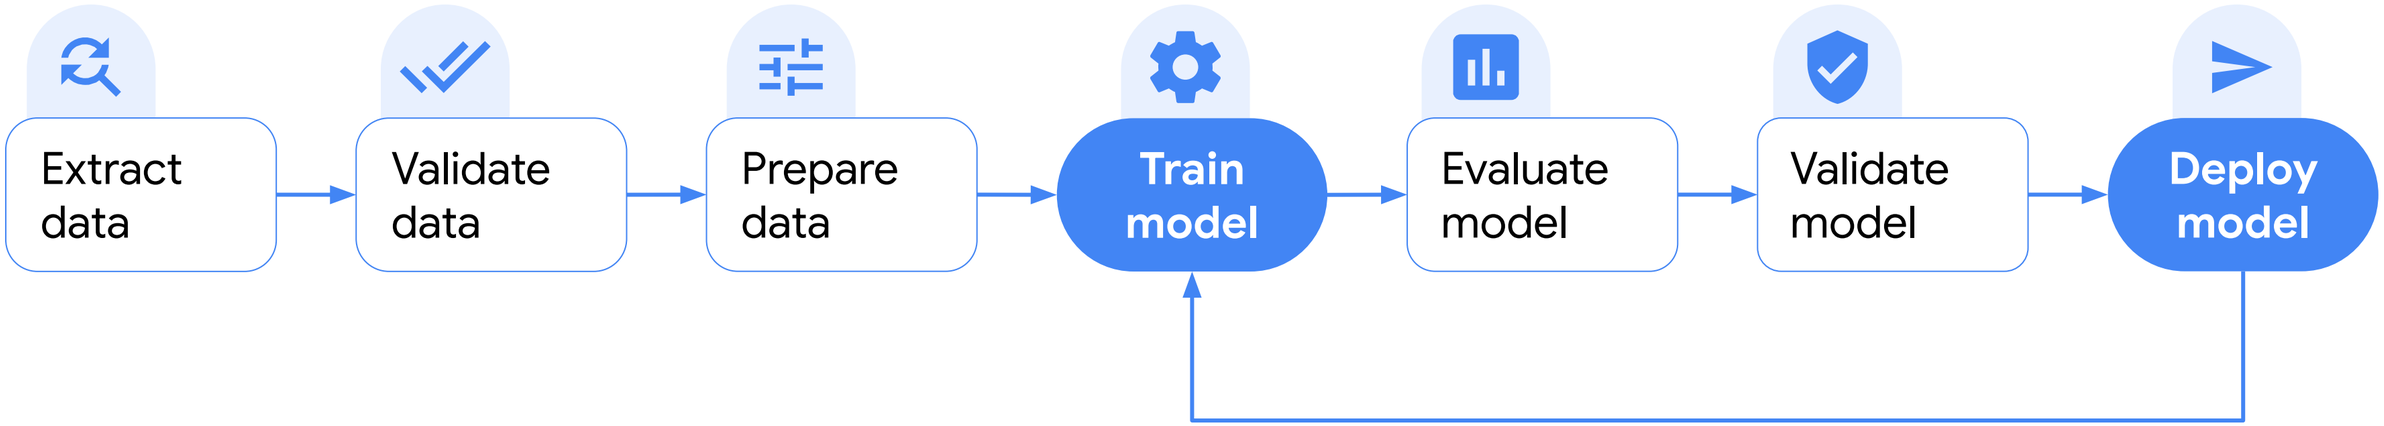

Pipelines are the backbone of production ML systems

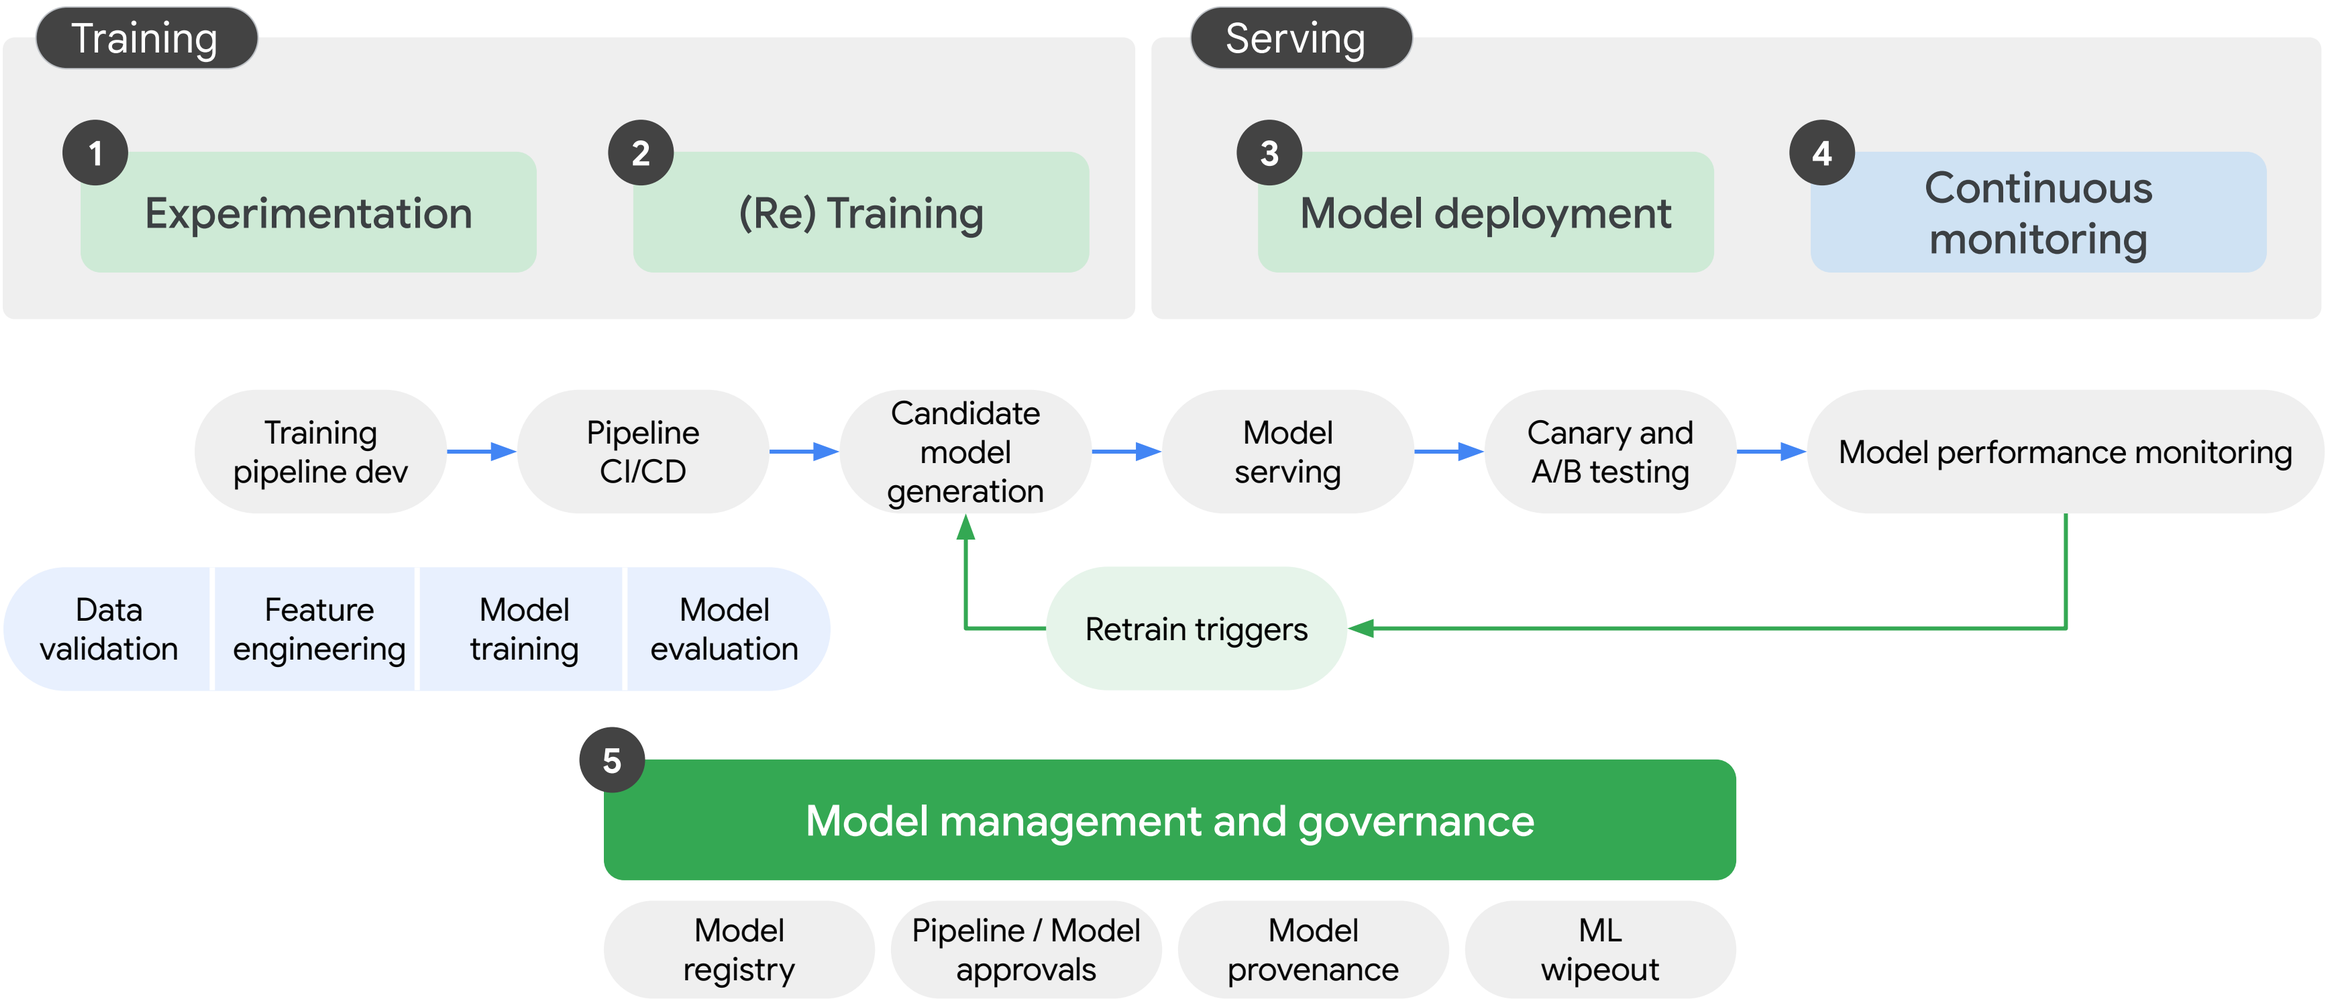

Pipelines product portfolio

Kubeflow

Kubeflow Pipelines

- Kubernetes-native.
- Open source.
- The industry standard for running ML Pipelines.

Google Cloud

AI Platform Pipelines - Hosted Beta

- Kubeflow pipelines running on Google Cloud.
- Optimized for GKE.
- Integrated with Google Cloud services.

Vertex Pipelines - Managed PREVIEW

- Fully managed and serverless.
- Allows users to focus on building their pipelines, scale easily, and pay only for the resources they use.

Write your pipeline

Easy to use Python SDKs

Build pipelines using Data Scientist friendly SDKs like TensorFlow Extended and Kubeflow Pipelines.

Rich, scalable pre-built components

We provide a rich set of pre-built components for common ML tasks, which leverage Google Cloud services.

```
@dsl.pipeline(pipeline_root=PIPELINE_ROOT, name="metadata-pipeline-v2")
def pipeline(message: str):
    importer = kfp.dsl.importer(
        artifact_uri="gs://ml-pipeline-playground/shakespeare1.txt",
        artifact_class=Dataset,
        reimport=False,
    )
    preprocess_task = preprocess(message=message)
    train_task = train(
        dataset_one=preprocess_task.outputs["output_dataset_one"],
        dataset_two=preprocess_task.outputs["output_dataset_two"],
        imported_dataset=importer.output,
        message=preprocess_task.outputs["output_parameter"],
        num_steps=5,
    )
    read_task = read_artifact_input(
        train_task.outputs["generic_artifact"]
    )
```

Key capabilities

Python SDKs

Data scientist friendly Python SDKs

1

Scalable

Run as many pipelines on as much data as you want.

2

Metadata and lineage

Store metadata for every artifact produced by the pipeline.

3

Monitoring UIs and APIs

Track and debug pipelines executions.

4

Security

Supports IAM, VPC-SC, and CMEK.

5

Cost-effective

Only pay for the pipelines you run and the resources they use.

6

Production ML Pipelines with Kubeflow

01	Ways to do ML on Google Cloud
02	Vertex AI Pipelines
03	AI Hub

AI Hub is a repository for AI assets

Don’t reinvent the wheel! Find and deploy ML pipelines.

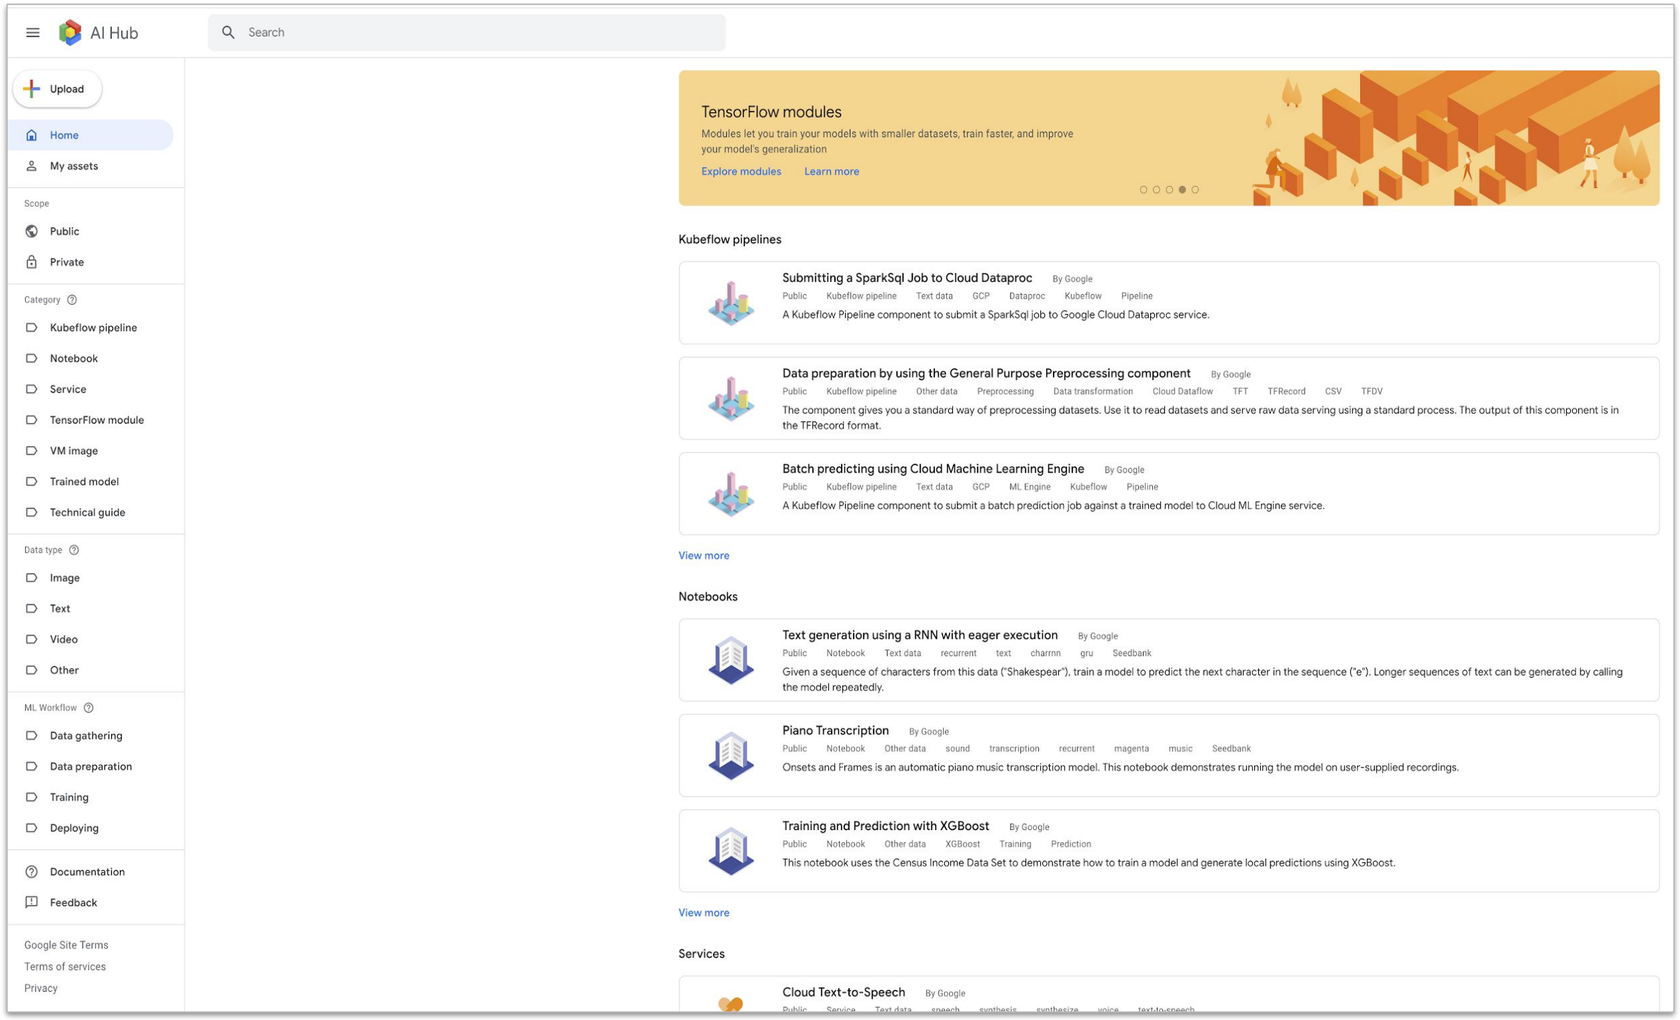

AI Hub stores various asset types

- Kubeflow pipelines and components
- Jupyter notebooks
- TensorFlow modules
- Trained models
- Services
- VM images

This is what a typical asset looks like

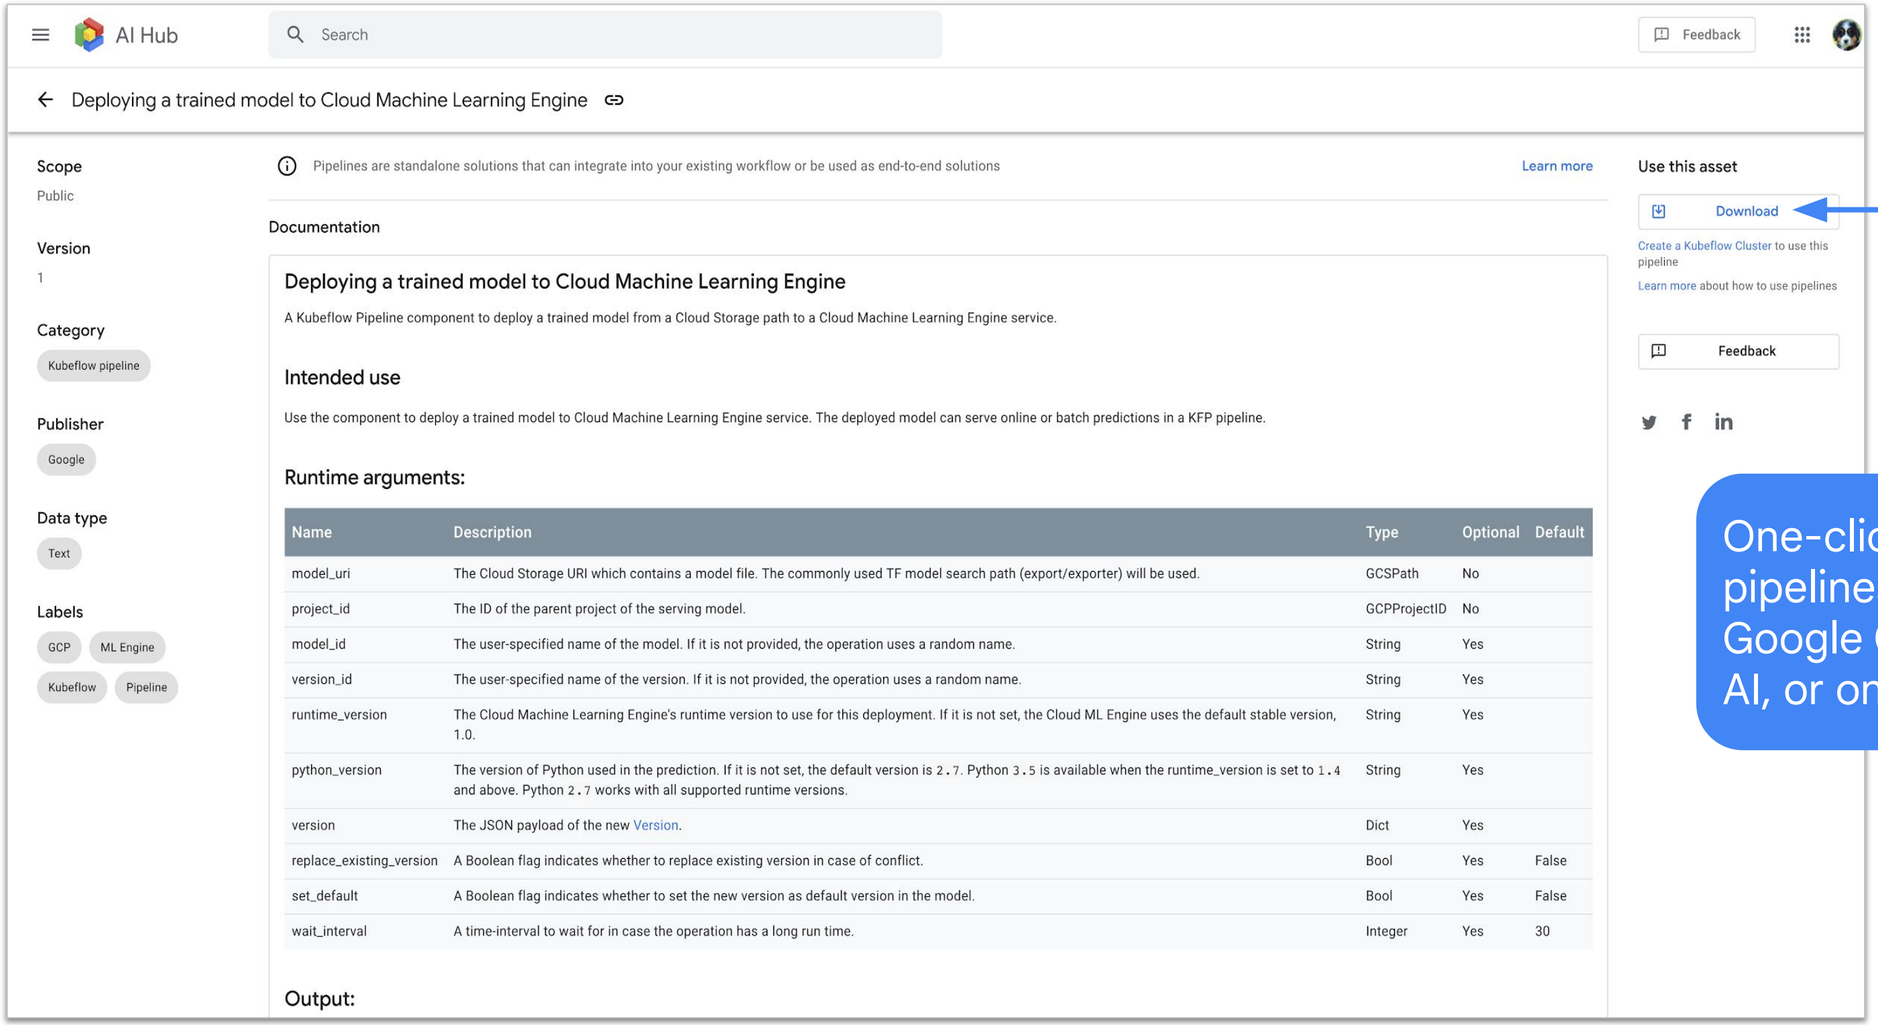

One-click deployment of ML pipelines via Kubeflow on Google Cloud as platform for AI, or on premise.

Assets on AI Hub are collected in two scopes: public assets and restricted assets

- Public scope are available to all AI Hub users.
- Restricted scope contains AI components that you have uploaded and assets that have been shared with you.

Lab Intro

Running Pipelines on Vertex AI

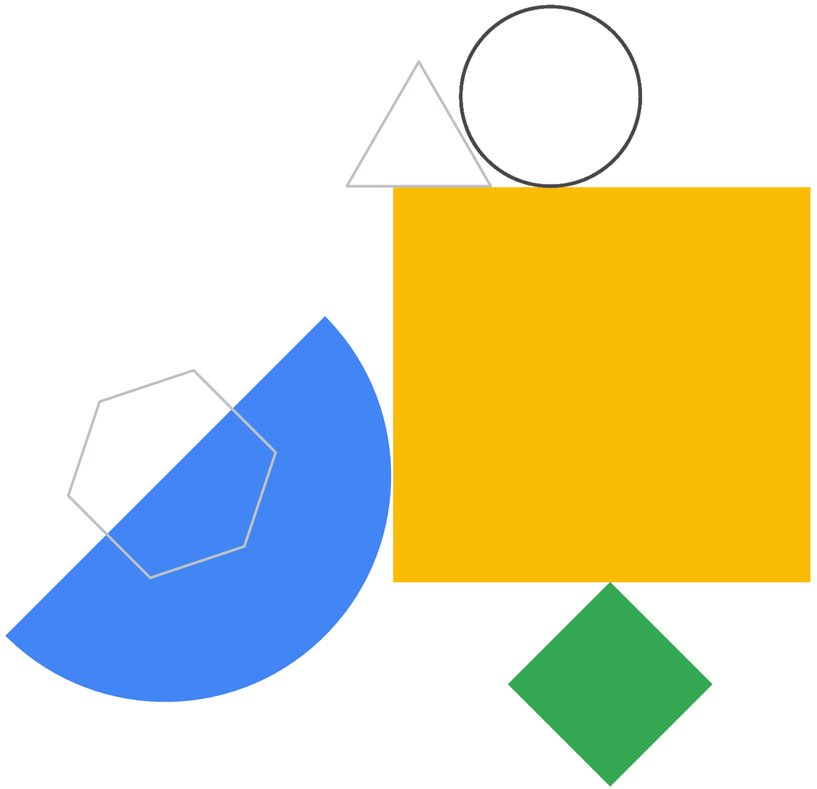

Lab objectives

01 Set up the project environment

02 Inspect and configure pipeline code

03 Execute the AI pipeline

Summary

- Use ML on Google Cloud using either:
  - Vertex AI (your model, your data)
  - AutoML (our models, your data)
- Use Vertex AI Pipelines to deploy end-to-end ML pipelines.
- Don’t reinvent the wheel for your ML pipeline! Leverage pipelines on AI Hub.