Google Cloud

Custom Model Building with SQL in BigQuery ML
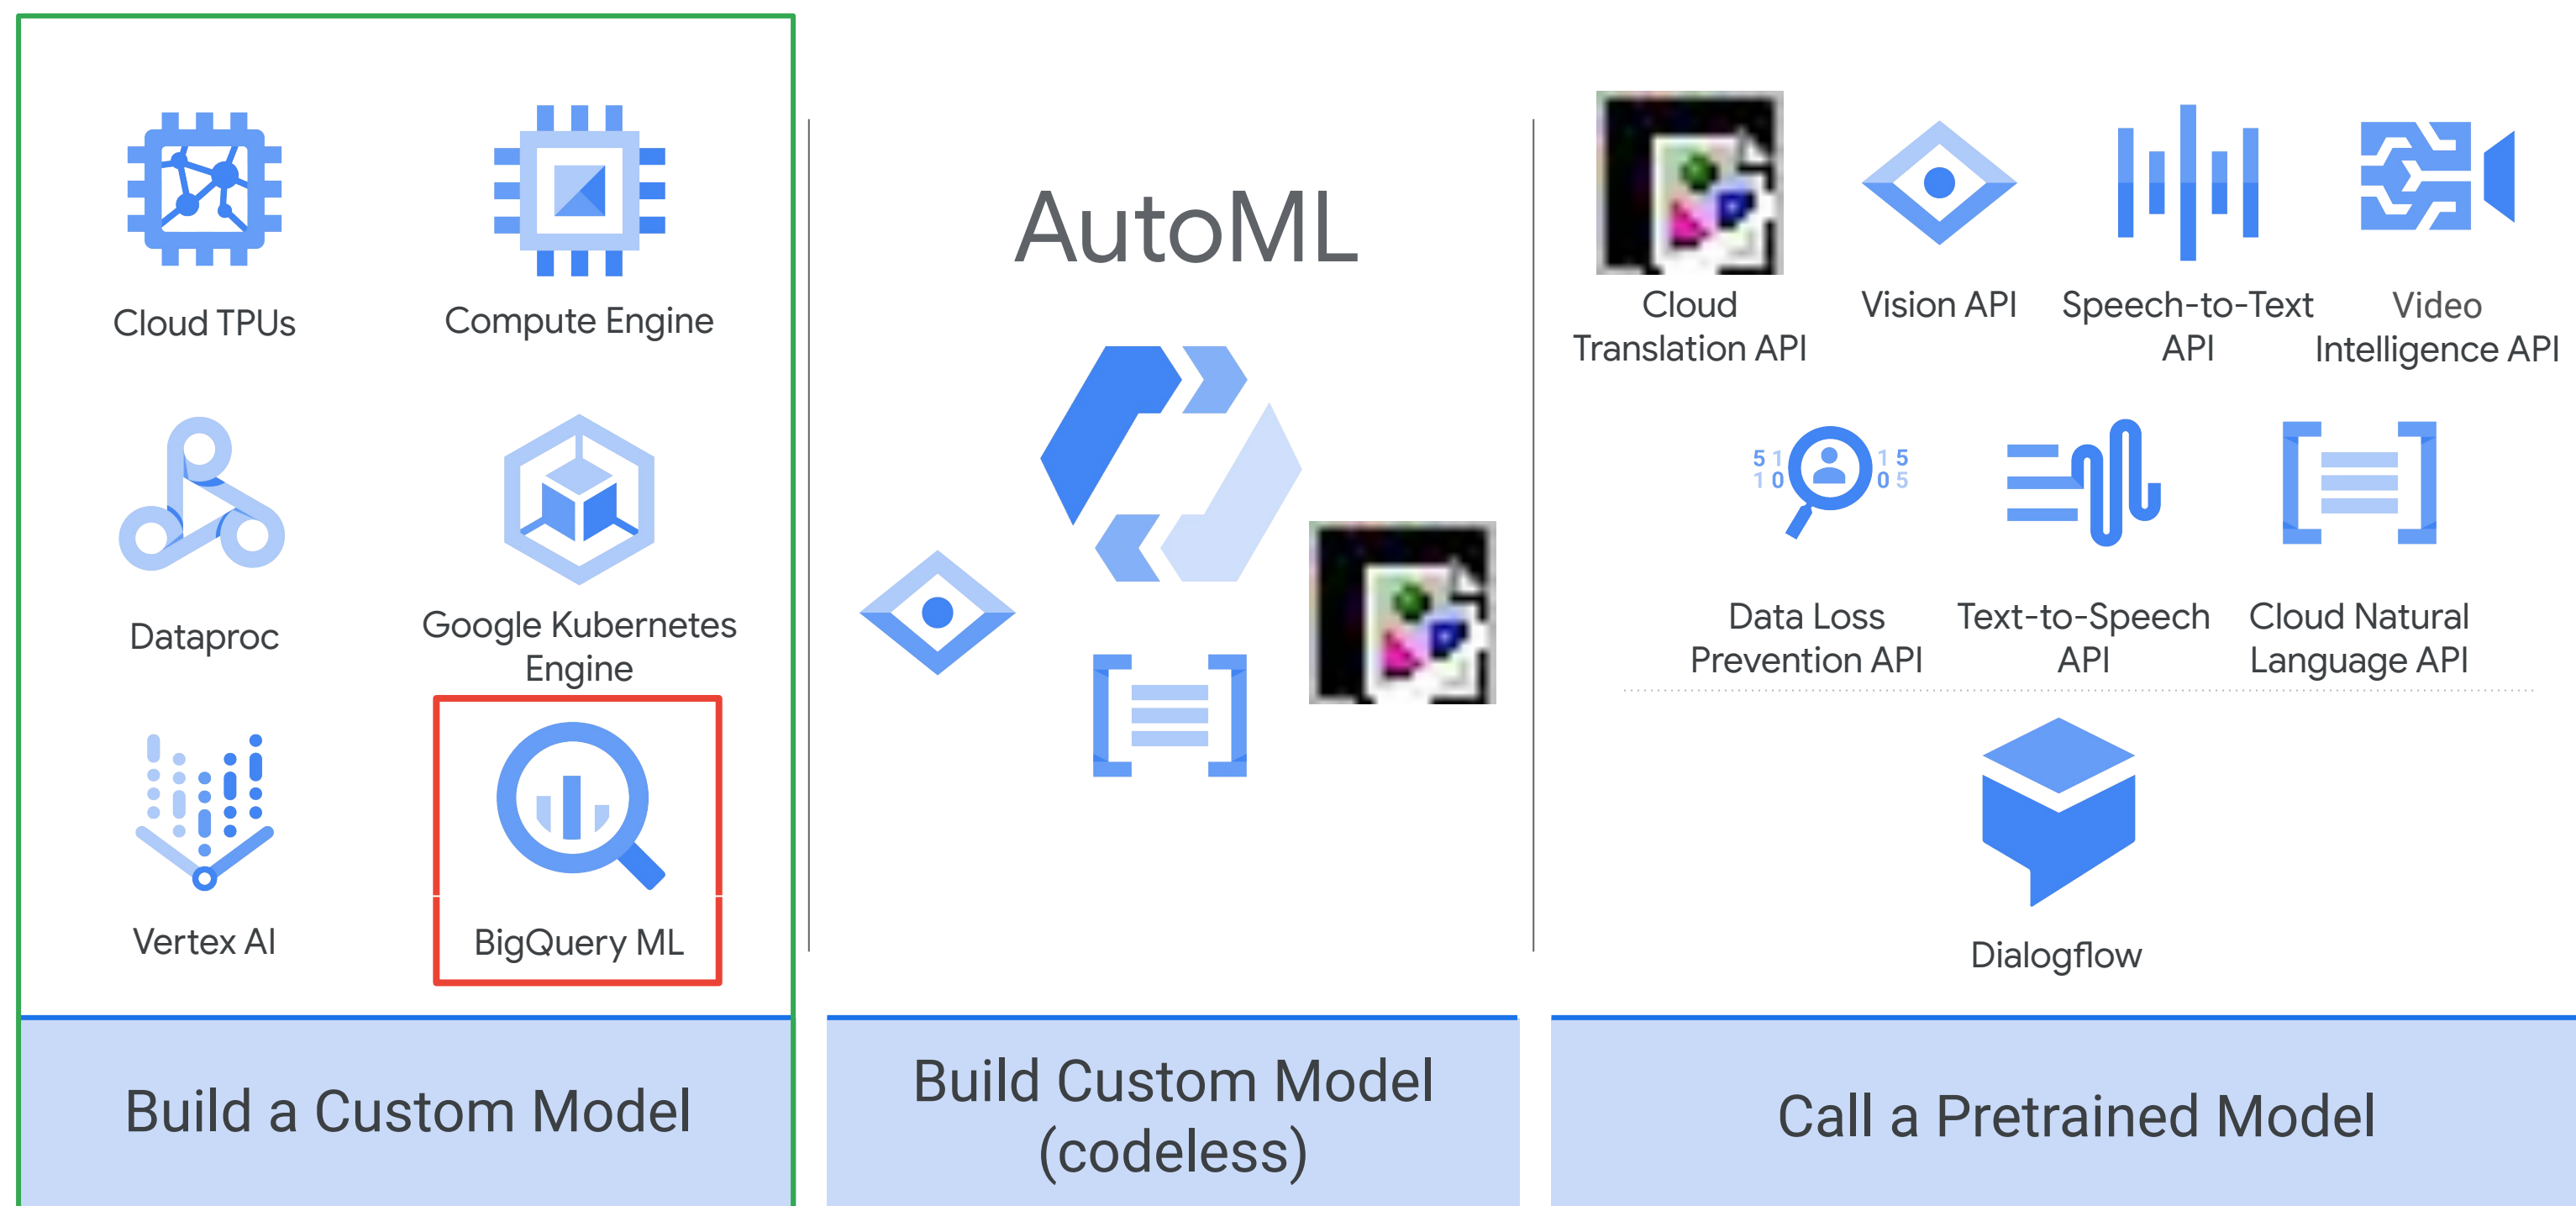
Agenda

BigQuery ML for Quick Model Building

Supported Models

Lab: Predict Bike Trip Duration with a Regression Model in BigQuery ML

Lab: Movie Recommendations in BigQuery ML

BigQuery ML is a way to build custom models

Cloud TPUs

Compute Engine

Dataproc

Google Kubernetes Engine

Vertex AI

BigQuery ML

AutoML

Cloud Translation API

Vision API

Speech-to-Text API

Video Intelligence API

Data Loss Prevention API

Text-to-Speech API

Cloud Natural Language API

Dialogflow

Build a Custom Model

Build Custom Model (codeless)

Call a Pretrained Model

Google Cloud

Working with BigQuery ML


```
FROM
    ML.EVALUATE(MODEL
`bqml_tutorial.sample_model`,
    TABLE eval_table)
```

① Dataset   ② Create/train   ③ Evaluate   ④ Predict/classify

```
CREATE MODEL `bqml_tutorial.sample_model`
OPTIONS(model_type='logistic_reg') AS
SELECT
```

```
FROM
    ML.PREDICT(MODEL
`bqml_tutorial.sample_model`,
    table game_to_predict) ) AS
predict
```

Google Cloud

Where was this article published?

1 Techcrunch

2 GitHub

3 NY Times

Unlikely Partnership in House Gives Lawmakers Hope for Border Deal

Representatives Nita M. Lowey and Kay Granger are the first women to lead the House Appropriations Committee. Their bond gives lawmakers optimism for the work to come.

By EMILY COCHRANE

Fitbit's newest fitness tracker is just for employees and health insurance members

Fitbit has a new fitness tracker, but it's one that you can't buy in stores. The company quietly uncorked the Inspire on Friday, releasing its first product that is available only to co...

1 hour ago   Jon Russell

Downloading the Android Studio Project Folder

FTC Engineering edited this page on Sep 19, 2017 · 1 revision

Downloading the Android Studio Project Folder

SQL query to extract data

```
SELECT
  url, title
FROM
  `bigquery-public-data.hacker_news.stories`
WHERE
  LENGTH(title) > 10
  AND LENGTH(url) > 0
LIMIT 10
```

*no clusters, no indexes, ad hoc query!

url	title
http://www.bbc.co.uk/news/business-27732743	Vodafone reveals direct government wiretaps
https://www.kickstarter.com/projects/appdocu/a...	Doc – App: The Human Story
http://www.starwebworld.com/android-jelly-bean...	Android Jelly Bean: Streaming Audio Through th...
http://www.myplanetdigital.com/digital_strateg...	Why Canadian Tech Entrepreneurs Need to Man/Wo...
http://startupislandconference.com/index.html	StartupConference June 13. - 16. 2013, HVAR Cr...
http://kopimism.org/	Kopimism Hactivism Meetup Tomorrow (Sunday) in...
http://unearthedgadget.com/xbox-live-gold-2/14...	Xbox Live Gold Membership Is It Really Worth -...
https://evertale.com	Evertale changes the way people remember
http://www.racketboy.com/retro/commodore-amiga...	Commodore Amiga: A Beginner's Guide
http://www.extremetech.com/extreme/156393-cold...	Cold fusion reactor "independently verified"

Use regex to get source + train on words of title

Create model

Query to extract training data

```
CREATE OR REPLACE MODEL advdata.txtclass
OPTIONS(model_type='logistic_reg', input_label_cols=['source'])
AS
```

```
WITH extracted AS (
...
)
, ds AS (
SELECT ARRAY_CONCAT(SPLIT(title, " "), ['NULL', 'NULL', 'NULL',
'NULL', 'NULL']) AS words, source FROM extracted
WHERE (source = 'github' OR source = 'nytimes' OR source =
'techcrunch')
)


SELECT
source,
words[OFFSET(0)] AS word1,
words[OFFSET(1)] AS word2,
words[OFFSET(2)] AS word3,
words[OFFSET(3)] AS word4,
words[OFFSET(4)] AS word5
FROM ds
```

Google Cloud

Evaluate model

```
SELECT * FROM ML.EVALUATE(MODEL advdata.txtclass)
```

precision	recall	accuracy	f1_score	log_loss	roc_auc
0.783	0.783	0.79	0.783	0.858	0.918

(BigQuery ML splits the training data and reports evaluation statistics on the held-out set)

Google Cloud

Predict using trained model

```
SELECT * FROM ML.PREDICT(MODEL advdata.txtclass,(
    SELECT 'government' AS word1, 'shutdown' AS word2, 'leaves' AS word3,
'workers' AS word4, 'reeling' AS word5
    UNION ALL SELECT 'unlikely', 'partnership', 'in', 'house', 'gives'
    UNION ALL SELECT 'fitbit', 's', 'fitness', 'tracker', 'is'
    UNION ALL SELECT 'downloading', 'the', 'android', 'studio', 'project'
))
```

"Batch prediction"

Row	predicted_source	word1	word2	word3	word4	word5
1	nytimes	government	shutdown	leaves	workers	reeling
2	nytimes	unlikely	partnership	in	house	gives
3	techcrunch	fitbit	s	fitness	tracker	is
4	techcrunch	downloading	the	android	studio	project

Google Cloud

Demo

Train a model with BigQuery ML to
predict NYC taxi fares

Agenda

BigQuery ML for Quick
Model Building

Supported Models

Lab: Predict Bike Trip Duration with
a Regression Model in BigQuery ML

Lab: Movie Recommendations in
BigQuery ML

Linear Classifier (Logistic regression)

```sql
#standardsql
CREATE OR REPLACE MODEL flights.ontime
OPTIONS
    (model_type='logistic_reg', input_label_cols=['on_time']) AS
SELECT
    IF(arr_delay < 15, 1, 0) AS on_time,
    carrier,
    origin,
    dest,
    dep_delay,
    taxi_out,
    distance
FROM
    `cloud-training-demos.flights.tzcorr`
WHERE
    arr_delay IS NOT NULL
```

DNN Classifier

```
#standardsql
CREATE OR REPLACE MODEL flights.ontime
OPTIONS
   (model_type='dnn_classifier', hidden_units = [47,29,18],
   input_label_cols=['on_time']) AS
SELECT
   IF(arr_delay < 15, 1, 0) AS on_time,
   carrier,
   origin,
   dest,
   dep_delay,
   taxi_out,
   distance
FROM
   `cloud-training-demos.flights.tzcorr`
WHERE
   arr_delay IS NOT NULL
```

xgboost Classifier

```
#standardsql
CREATE OR REPLACE MODEL flights.ontime
OPTIONS
    (model_type='boosted_tree_classifier', input_label_cols=['on_time']) AS
SELECT
    IF(arr_delay < 15, 1, 0) AS on_time,
    carrier,
    origin,
    dest,
    dep_delay,
    taxi_out,
    distance
FROM
    `cloud-training-demos.flights.tzcorr`
WHERE
    arr_delay IS NOT NULL
```

Linear Regression

```
CREATE OR REPLACE MODEL
    taxi.taxifare_dnn OPTIONS (model_type='linear_reg', labels=['fare_amount']) AS
SELECT
    fare_amount,
    hourofday, dayofweek,
    pickuplon, pickuplat, dropofflon, dropofflat,
    passenger_count
FROM
    `taxi.taxi3m`
```

DNN Regression

```
CREATE OR REPLACE MODEL
    taxi.taxifare_dnn OPTIONS (model_type='dnn_regressor',
    hidden_units=[144,89,55], labels=['fare_amount']) AS
SELECT
    fare_amount,
    hourofday, dayofweek,
    pickuplon, pickuplat, dropofflon, dropofflat,
    passenger_count
FROM
    `taxi.taxi3m`
```

xgboost Regression

```
CREATE OR REPLACE MODEL
    taxi.taxifare_xgboost
    OPTIONS (model_type='boosted_tree_regressor', labels=['fare_amount']) AS
SELECT
    fare_amount,
    hourofday, dayofweek,
    pickuplon, pickuplat, dropofflon, dropofflat,
    passenger_count
FROM
    `taxi.taxi3m`
```

Google Cloud

Train on TF, predict with BigQuery

```
CREATE OR REPLACE MODEL advdata.txtclass_tf2
OPTIONS (model_type='tensorflow',

model_path='gs://cloud-training-demos-ml/txtcls/trained_finetune_native/export/exporter/1
549825580/*')
```

```
SELECT
  input,
  (SELECT AS STRUCT(p, ['github', 'nytimes', 'techcrunch'][ORDINAL(s)]) prediction FROM
    (SELECT p, ROW_NUMBER() OVER() AS s FROM
      (SELECT * FROM UNNEST(dense_1) AS p))
  ORDER BY p DESC LIMIT 1).*

FROM ML.PREDICT(MODEL advdata.txtclass_tf2,
(
SELECT 'Unlikely Partnership in House Gives Lawmakers Hope for Border Deal' AS input
UNION ALL SELECT "Fitbit\'s newest fitness tracker is just for employees and health
insurance members"
UNION ALL SELECT "Show HN: Hello, a CLI tool for managing social media"
))
```

Recommendation engine (matrix factorization)

```
create or replace model models.suggested_products_1or2_example
options(model_type='matrix_factorization',
        user_col='user_id', item_col='product_id', rating_col='rating',
        l2_reg=10)
AS

with purchases AS (
 select product_id, user_id from
 operations.orders_with_lines, unnest(order_lines)
),

total_purchases as (
select product_id, user_id, count(*) as numtimes
from purchases
group by product_id, user_id
)

select
product_id, user_id,
IF(numtimes < 2, 1, 2) AS rating
FROM total_purchases
```

So what do we recommend for a given set of users?

```
with users AS (
SELECT
user_id, count(*) as num_orders
from operations.orders_with_lines
group by user_id
order by num_orders desc
limit 10
),

products as (
select product_id, count(*) as num_orders
from operations.orders_with_lines, unnest(order_lines)
group by product_id
order by num_orders desc
limit 10
)

SELECT * FROM ML.PREDICT(MODEL models.suggested_products_1or2,
(SELECT user_id, product_id
FROM users, products)
)
```

Google Cloud

So what do we recommend for a given set of users?

Row	predicted_rating	user_id	product_id
1	1.5746015507788755	101797	26209
2	1.8070705987455633	101797	13176
3	1.7171094544245578	101797	27845
4	1.9763373899260837	101797	47209
5	1.8659380090171271	101797	21137
6	1.721610848530093	101797	47766
7	1.9516130703939483	101797	21903

Clustering

```
CREATE OR REPLACE MODEL
demos_eu.london_station_clusters
OPTIONS(model_type='kmeans', num_clusters=4,
standardize_features = true) AS


WITH hs AS …,
stationstats AS …

SELECT * except(station_name, isweekday)
from stationstats
```

1. 4 clusters (hardcoded)

2. Standardize features since different dynamic ranges

3. Remove the cluster "id" fields (keep just the attributes)

Google Cloud

Which cluster?

```
WITH hs AS ...,
stationstats AS ...,

SELECT * except(nearest_centroids_distance)
FROM ML.PREDICT(MODEL demos_eu.london_station_clusters,
(SELECT * FROM stationstats WHERE
REGEXP_CONTAINS(station_name, 'Kennington')))
```

| Row	CENTROID_ID	station_name	isweekday	duration	num_trips	bikes_count	distance_from_city_center
1	3	Kennington Lane Tesco, Vauxhall	weekday	911.5810637908974	5471	9	1.8345619962343163
2	3	Kennington Lane Rail Bridge, Vauxhall	weekday	979.3919952622995	20263	19	2.175032834765301
3	4	Doddington Grove, Kennington	weekday	1397.7189755200225	7067	28	1.468140527379382
4	4	Kennington Cross, Kennington	weekday	911.5238777770538	15349	35	1.4625875338501981

Find cluster attributes

```
WITH T AS (
SELECT
centroid_id,
ARRAY_AGG(STRUCT(numerical_feature AS name, ROUND(feature_value,1) AS value) ORDER BY
centroid_id) AS cluster
FROM ML.CENTROIDS(MODEL demos_eu.london_station_clusters)
GROUP BY centroid_id
)
SELECT
CONCAT('Cluster#', CAST(centroid_id AS STRING)) AS centroid,
(SELECT value from unnest(cluster) WHERE name = 'duration') AS duration,
(SELECT value from unnest(cluster) WHERE name = 'num_trips') AS num_trips,
(SELECT value from unnest(cluster) WHERE name = 'bikes_count') AS bikes_count,
(SELECT value from unnest(cluster) WHERE name = 'distance_from_city_center') AS
distance_from_city_center
FROM T
ORDER BY centroid_id ASC
```

Google Cloud

Visualize attributes in Google Data Studio

Row	centroid	duration	num_trips	bikes_count	distance_from_city_center
1	Cluster#1	3079.5	3026.1	14.0	6.2
2	Cluster#2	1564.0	3635.1	11.5	6.5
3	Cluster#3	1319.6	9654.8	6.4	2.9
4	Cluster#4	1527.7	4846.8	22.6	3.5

centroid ▲	bikes_count	distance...	duration	num_trips
Cluster#1				
Cluster#2				
Cluster#3				
Cluster#4				

Google Cloud

Use the transform clause
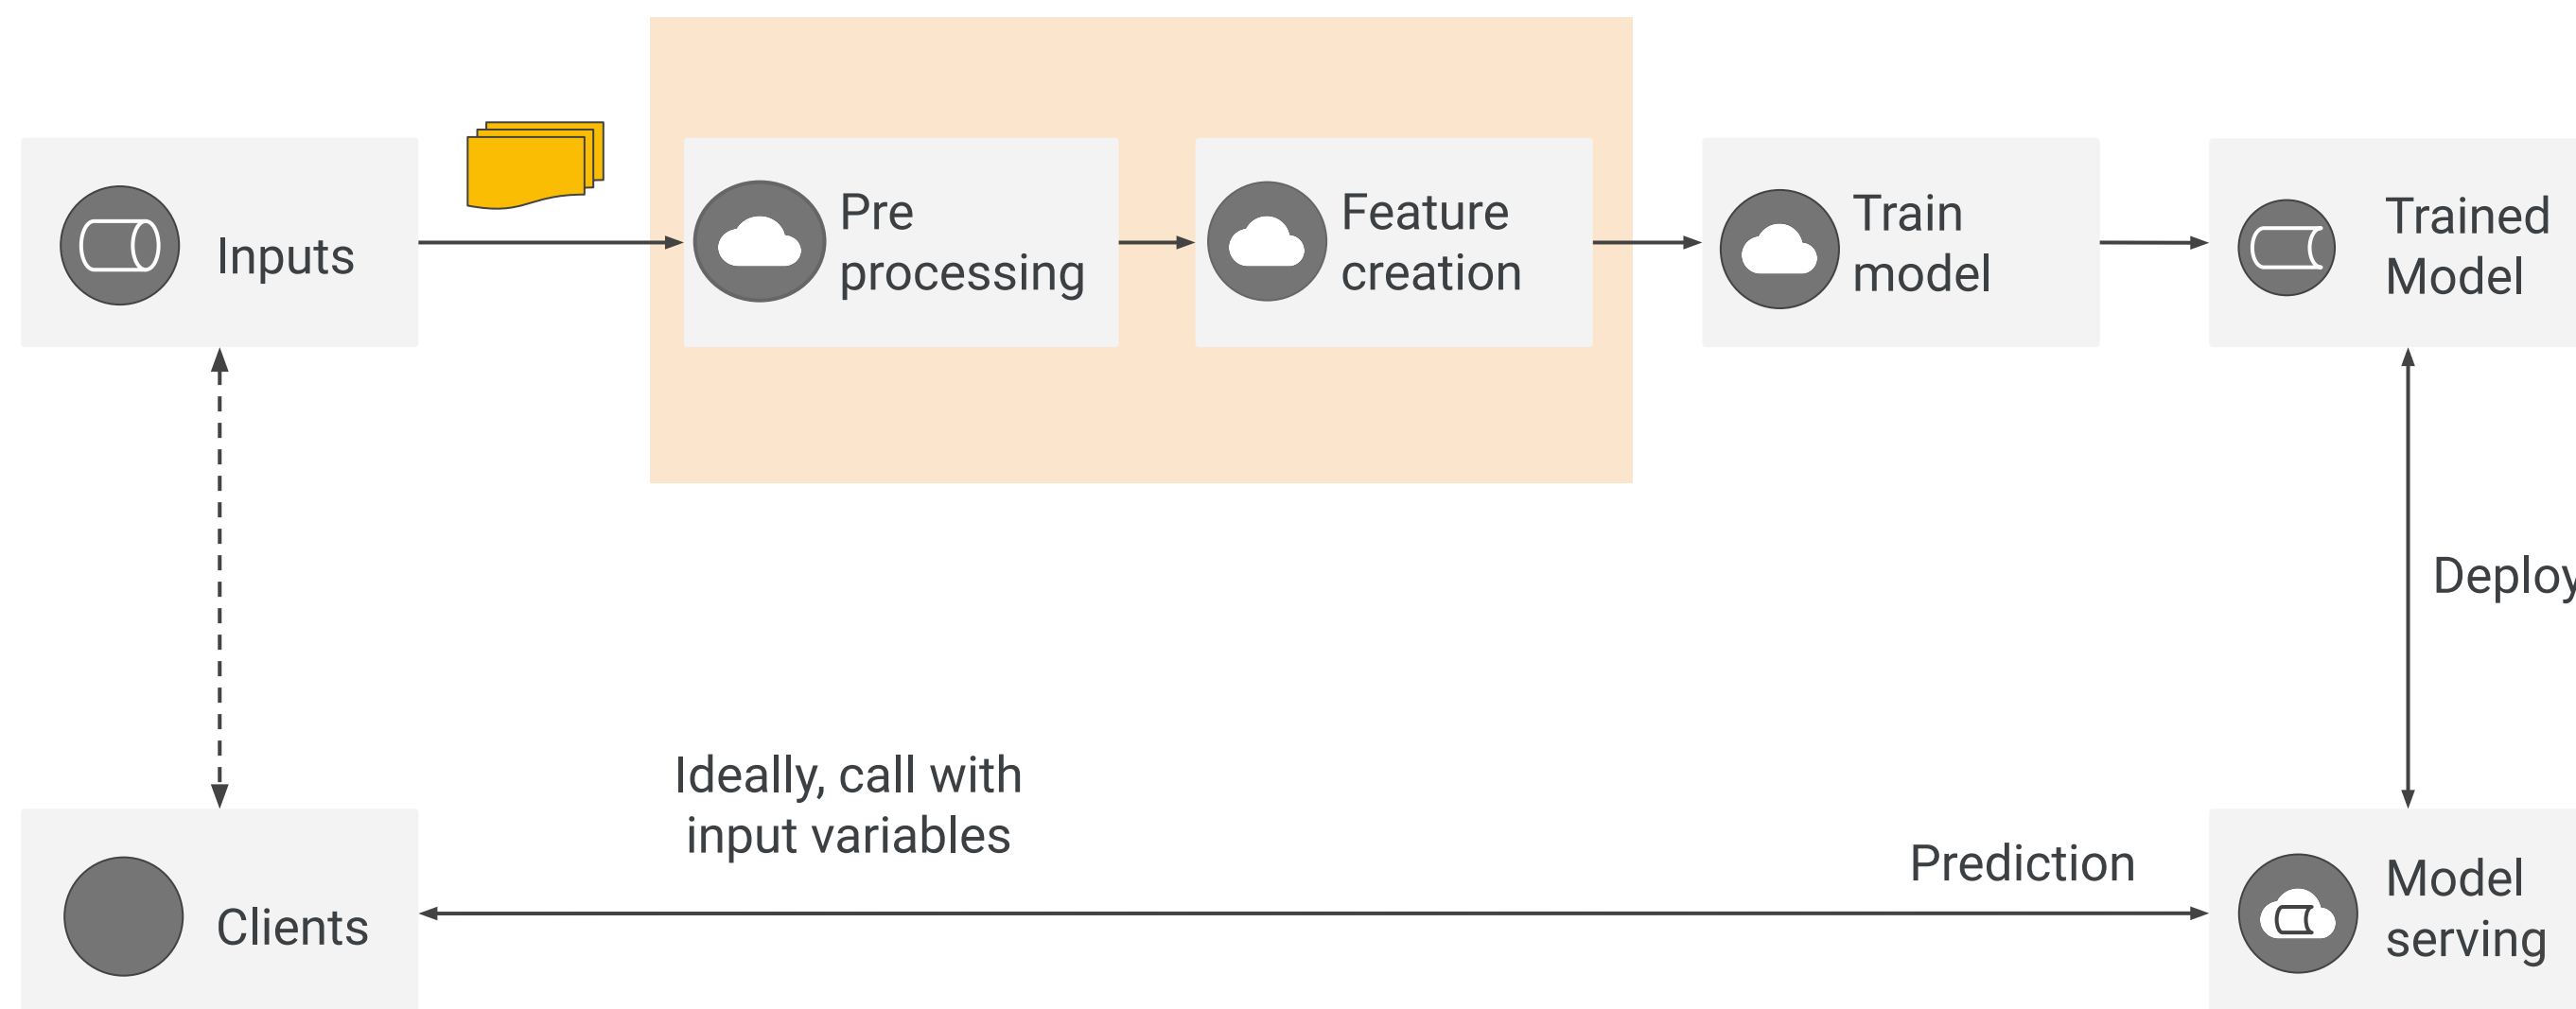
TRANSFORM ensures transformations are automatically applied during ML.PREDICT

```
CREATE OR REPLACE MODEL ch09edu.bicycle_model
OPTIONS(input_label_cols=['duration'],
        model_type='linear_reg')
AS

SELECT
  duration
  , start_station_name
  , CAST(EXTRACT(dayofweek from start_date) AS STRING)
        as dayofweek
  , CAST(EXTRACT(hour from start_date) AS STRING)
        as hourofday
FROM
  `bigquery-public-data.london_bicycles.cycle_hire`
```

```
SELECT * FROM ML.PREDICT(MODEL ch09edu.bicycle_model,(
  350 AS duration
  , 'Kings Cross' AS start_station_name
  , '3' as dayofweek
  , '18' as hourofday
))
```

```
CREATE OR REPLACE MODEL ch09edu.bicycle_model
OPTIONS(input_label_cols=['duration'],
        model_type='linear_reg')
TRANSFORM(
  SELECT * EXCEPT(start_date)
  , CAST(EXTRACT(dayofweek from start_date) AS STRING)
        as dayofweek
  , CAST(EXTRACT(hour from start_date) AS STRING)
        as hourofday
)
AS
SELECT
  duration, start_station_name, start_date
FROM
  `bigquery-public-data.london_bicycles.cycle_hire`
```

```
SELECT * FROM ML.PREDICT(MODEL ch09edu.bicycle_model,(
  350 AS duration
  , 'Kings Cross' AS start_station_name
  , CURRENT_TIMESTAMP() as start_date
))
```

Google Cloud

Reminder: BigQuery ML Cheatsheet

- **Label** = `alias a column as 'label' or specify column in OPTIONS using input_label_cols`

- **Feature** = passed through to the model as part of your SQL SELECT statement
  `SELECT * FROM ML.FEATURE_INFO(MODEL `mydataset.mymodel`)`

- **Model** = an object created in BigQuery that resides in your BigQuery dataset

- **Model Types** = Linear Regression, Logistic Regression
  `CREATE OR REPLACE MODEL <dataset>.<name>`
  `OPTIONS(model_type='<type>') AS`
  `<training dataset>`

- **Training Progress** = `SELECT * FROM ML.TRAINING_INFO(MODEL `mydataset.mymodel`)`

- **Inspect Weights** = `SELECT * FROM ML.WEIGHTS(MODEL `mydataset.mymodel`, (<query>))`

- **Evaluation** = `SELECT * FROM ML.EVALUATE(MODEL `mydataset.mymodel`)`

- **Prediction** = `SELECT * FROM ML.PREDICT(MODEL `mydataset.mymodel`, (<query>))`

Agenda

BigQuery ML for Quick
Model Building

Supported Models

Lab: Predict Bike Trip Duration with
a Regression Model in BigQuery ML

Lab: Movie Recommendations in
BigQuery ML

Google Cloud

Lab Intro

Predict Bike Trip Duration with a Regression Model in BigQuery ML

Agenda

BigQuery ML for Quick
Model Building

Supported Models

Lab: Predict Bike Trip Duration with
a Regression Model in BigQuery ML

Lab: Movie Recommendations in
BigQuery ML

Google Cloud

Lab Intro

Movie Recommendations in
BigQuery ML

Module summary

You can train and evaluate machine learning models directly in BigQuery.

Google Cloud